

CS 5/7320
Artificial Intelligence

Learning from Examples:
Learning to Score a Tic-
Tac-Toe board

AIMA Chapter 19

Michael Hahsler

Discussion


Intro: Learning from Examples

What can be learned?

Learning from Examples

Up until now in this course:

- **Hand-craft algorithms** to make rational/optimal or at least good decisions.
Examples: Search strategies, heuristics, and constructing Bayesian networks.




Issues

- We may not be able to anticipate all possible future scenarios.
- We may have examples, but we do not know how to implement a solution.


Supervised Machine Learning

- Uses observations: training data with the correct answers.
- Learn a function (model) to map an input (e.g., state) to an output (e.g., action) representing the desired behavior.
- Examples:
 - Use a naïve Bayesian classifier to distinguish between spam/non-spam.
 - Learn a playout policy to simulate games (current board -> good move)



Learning Components of an Agent

- We can learn many different components of an agent from examples
- **Example:** Learning components of a model-based reflex agent



Typical Use of Supervised ML for Intelligent Agents


Learn a Policy	Learn Evaluation Functions	Learn Perception and Actuation	Compressing Tables
<ul style="list-style-type: none">Classification: Directly learn the best action for each state from examples. $a = h(state \ features)$This model can also be used as a playout policy for Monte Carlo tree search with data from self-play.	<ul style="list-style-type: none">Regression: Learn evaluation functions to estimate state utilities. $eval = h(state \ features)$Can learn a heuristic for heuristic alpha-beta search.For reinforcement learning we can learn action values $q(state, action)$.	<ul style="list-style-type: none">Natural language processing: Use deep learning / word embeddings / language models to understand concepts, translate between languages, or generate text.Speech recognition: Identify the most likely sequence of words.Vision: Object recognition in images/videos. Generate images/video.Robotics: Learn how to move safely.	<ul style="list-style-type: none">Neural networks can be used as a compact representation of tables that do not fit in memory. E.g.,<ul style="list-style-type: none">Joint and conditional probability tablesState utility tables (i.e., an evaluation function)Q-Value tables in reinforcement learning

Bottom line: Learning a function is often more effective than hard-coding it. However, we do not always know how it performs for rare and edge cases!

Example: Tic-Tac-Toe

Approaches to implement/learn an evaluation function.

Utility of Being in a State: State Value



Approach: Try all actions and pick the action that leads to the state with the largest state value.

$$\operatorname{argmax}_{a \in \text{actions}(s)} \text{value}(\text{result}(s, a))$$

The state value is defined as:

$$\text{value}(s) = \mathbb{E}_{\pi_x, \pi_o}[U(s)],$$

with the playout policies π_x, π_o of player x and o.



Expectation is used for
stochastic playout
strategies!

Questions:

- What playout policies do we use? Random, optimal, ...
- How do we estimate the expected value?

Minimax and Variations

$$value(s) \approx \mathbb{E}_{\pi_x, \pi_o}[U(s)]$$

State s	$\text{minimax}(s)$	Expert code $\text{eval}(s)$		
	1			
	0			
:				

Playout policy
 π_x, π_o

optimal play

Method

Tree search

Expert Evaluation Function

$$value(s) \approx \mathbb{E}_{\pi_x, \pi_o}[U(s)]$$

State s	$\text{minimax}(s)$	Expert code $eval(s)$	Monte Carlo simulation $eval(s)$	
	1	+0.4		
	0	-0.4		
:				

Expert Heuristic

.4 for each winning opportunity
.4 for each loosing opportunity

Playout policy
 π_x, π_o

optimal play

none

Method

Tree search

Code + cutoff
search

Monte Carlo Simulation

$$value(s) \approx \mathbb{E}_{\pi_x, \pi_o}[U(s)]$$

Monte Carlo Simulation
 $\approx \frac{1}{N} \sum_{i=1}^N u_i(s)$

State s	$minimax(s)$	Expert code $eval(s)$	Monte Carlo simulation $eval(s)$	Machine Learning $h(s)$
	1	+0.4	0.634	
	0	-0.4	-0.021	
:				

Playout policy
 π_x, π_o

optimal play

none

Random or
other

Method



Tree search

Code + cutoff
search

Simulation

Machine Learning

$$value(s) \approx E_{\pi_x, \pi_o}[U(s)]$$

State s	$\text{minimax}(s)$	Expert code $eval(s)$	Monte Carlo simulation $eval(s)$	Machine Learning $h(s)$
	1	+0.4	0.634	
	0	-0.4	-0.021	
:				

Playout policy
 π_x, π_o

optimal play

none

Random or
other

Method



Tree search

Code + cutoff
search

Simulation

Machine Learning

$$value(s) \approx \mathbb{E}_{\pi_x, \pi_o}[U(s)]$$

State s	$\text{minimax}(s)$	Expert code $\text{eval}(s)$	Monte Carlo simulation $\text{eval}(s)$	Machine Learning $h(s)$
	1	+0.4	0.634	0.325
	0	-0.4	-0.021	0.042
:				

train ML

Playout policy
 π_x, π_o

optimal play

none

Random or
other

For training
data

Method

Tree search

Code + cutoff
search

Simulation

Machine
learning

All methods
hopefully produce a
consistent state
order.

Reinforcement Learning (RL)

Tries to learn state values directly from rewards resulting from

- **Model-based RL:** Uses a complete probabilistic model of the environment.
- **Model-free RL:** Learns from interactions with the environment.
 - Online learning
 - Offline learning (can use supervised learning from a sampled dataset)

Machine learning for an evaluation function can be seen as a special case of model-free offline RL called (neural) fitted Q-learning with replay memory.

We will talk about other RL methods later.